

5. Security Controls

Prof. Dr. Marc Rennhard

Institut für angewandte Informationstechnologie InIT

ZHAW School of Engineering

rema@zhaw.ch

Content

This chapter provides an **overview of security controls** that are likely to be relevant when developing secure applications; most of them were discussed in detail in the course Internet-Sicherheit (ISI):

- Cryptographic primitives
 - User authentication
 - Secure communication protocols
 - Authorization
 - Firewalls
 - Intrusion detection / prevention systems
-
- For additional information refer to the corresponding ISI course material


Goals

- You have an **overview of a variety of security controls** that are relevant when developing secure (distributed) systems
- You can describe the **functionality** of each security control
- You know the **application areas** of the different security controls
- You know the **strengths and limitations** of each security control

Cryptographic Primitives

Secret Key Ciphers (1)

- Secret key ciphers serve to **encrypt data to achieve confidentiality**



- Popular ciphers: **AES** (block cipher), **RC4** (stream cipher)
- The **key must be secretly exchanged** beforehand
- The key length should be at least **128 bits**
- As always in cryptography, it is extremely important that the key material is picked "**as randomly as possible**"

Secret Key Ciphers (2)

- Applications:
 - As a building block in **secure communication protocols** to protect “data in motion”
 - To **protect “data at rest”**, e.g. disk encryption
- Strengths:
 - **Strong and widely distributed standards** (e.g. AES) available
 - **Very fast**, even on small devices feasible
- Limitations:
 - Only provide encryption but **do not solve the key exchange problem**
 - When used for “raw encryption”, secret key ciphers only provide confidentiality, but attacks on integrity are still possible → combine with appropriate integrity-protection

Public Key Ciphers (1)

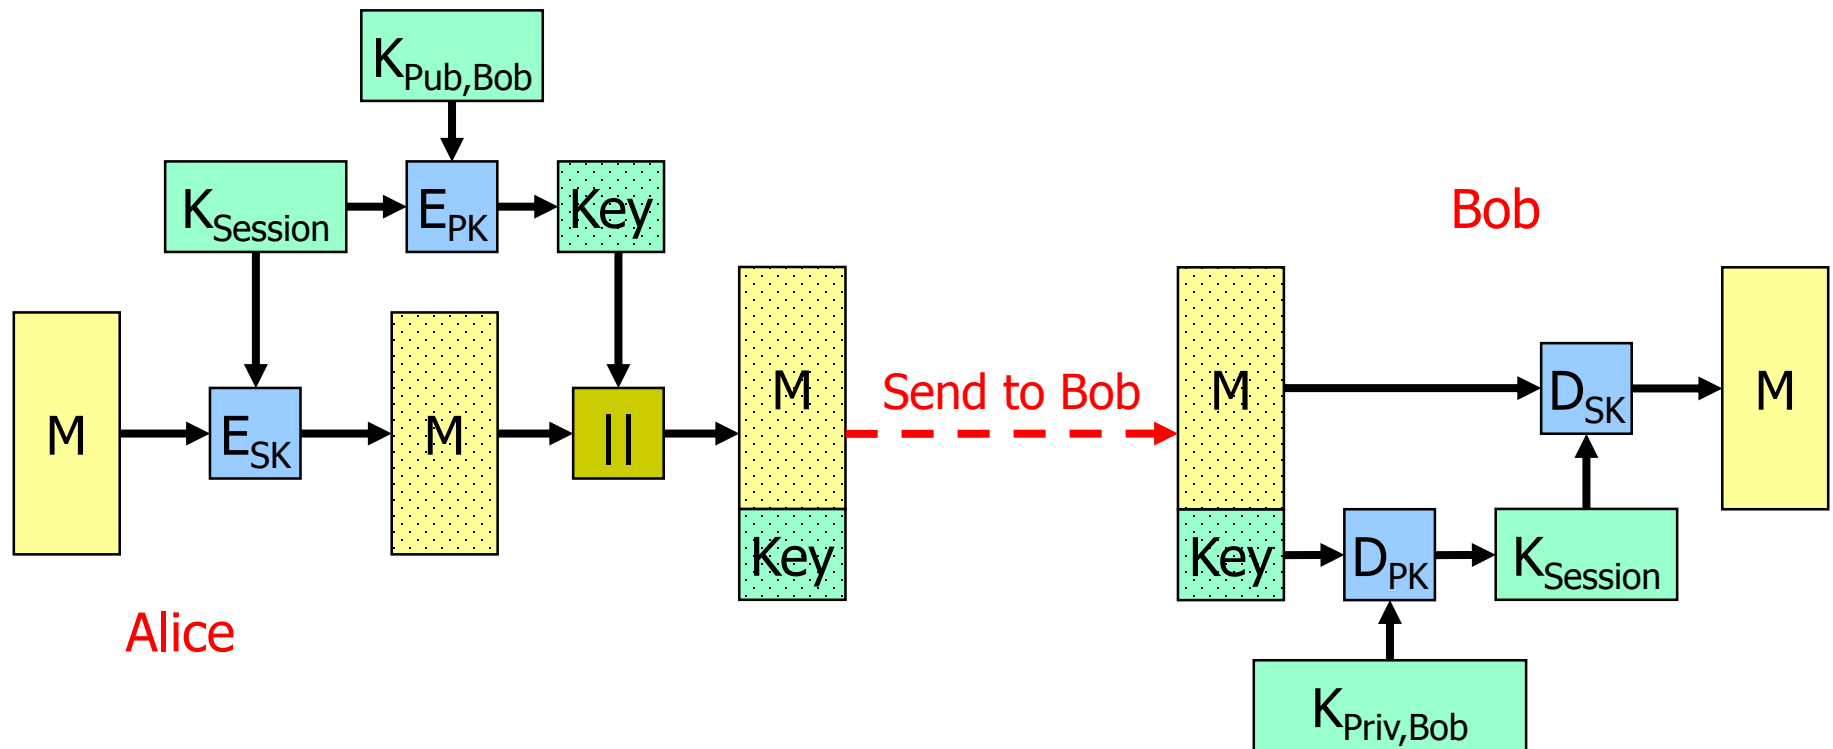
- In contrast to secret key cryptography, **two different keys** are used for encryption and decryption
 - To encrypt a message, the **public key** is used
 - To decrypt a message, the **private key** is used
 - Corresponding public and private key belong together: **key pair**
- A person Alice can **generate a key pair and publish her public key**
 - Other persons can **take Alice's public key and encrypt a message for her**
 - Only Alice can decrypt these messages using her own private key
 - Decryption of a message using the public key is not possible
- Popular public key algorithms:
 - **RSA** (use keys with a modulus size of at least 2'048 bits)
 - Less frequently used: Elliptic Curve Cryptography (**ECC**)

Public Key Ciphers (2)

- Applications:
 - As a building block in **secure communication protocols** to exchange a secret key
- Strengths:
 - **Strong and widely distributed standards** (e.g. RSA) available
 - **Solves the key exchange problem** of secret key cryptography (see hybrid encryption), assuming the authenticity of the public key can be verified
- Limitations:
 - **Much slower** than secret key cryptography, not suited for bulk data encryption

Hybrid Encryption

- Combining secret and public key cryptography: Alice has Bob's public key and wants to send him an encrypted (long) message
 - The message itself is encrypted using a secret key cipher with a **randomly selected session key**
 - The session key is encrypted with **Bob's public key**



Cryptographic One-Way Hash Functions (1)

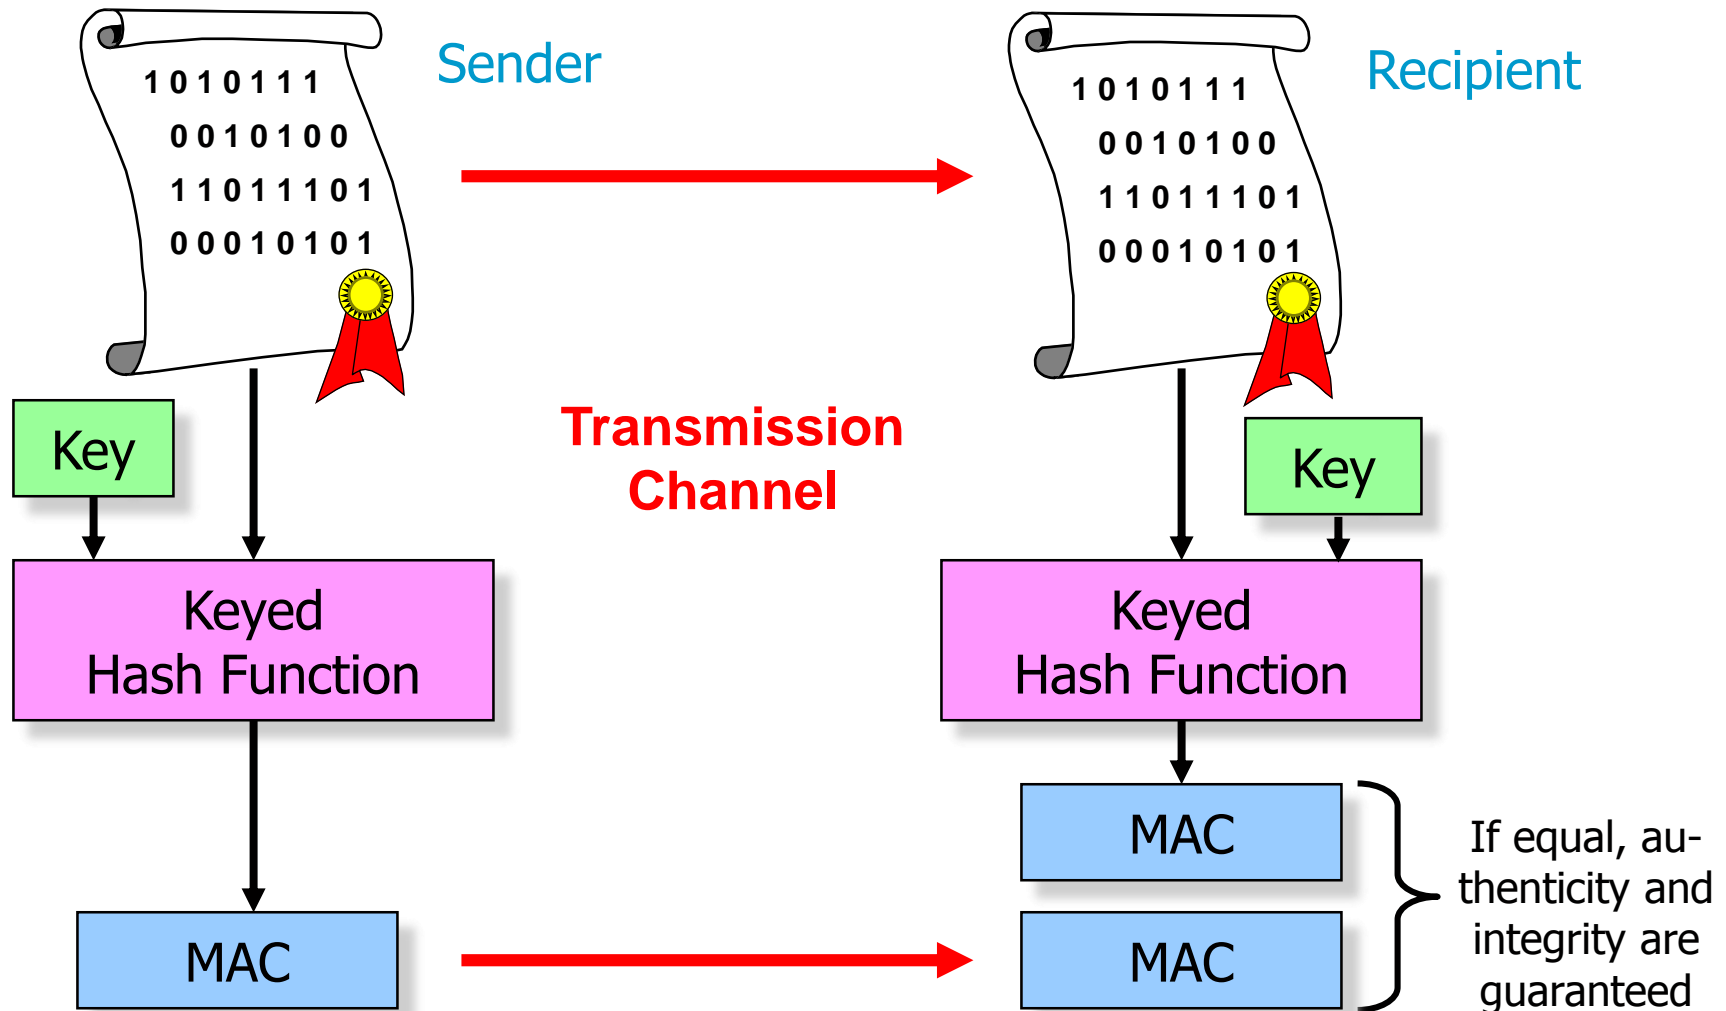
- A cryptographic one-way hash function maps a **variable-length input bit string** (the message) to a **fixed-sized output bit string** (the hash or message digest) → many-to-one mapping
- Important properties:
 - Given a hash, it is practically impossible to find a message that produces the hash (**one-way**)
 - It is practically impossible to find any two messages that map to the same hash (**collision-free**)
- All properties fulfilled: **in practice**, a message produces a unique hash and a hash belongs to a particular message → **considered as a one-to-one mapping**
- Typical hash lengths: **128, 160, 256, 512 bits**
- Popular hash functions: **MD5, SHA-1, (SHA-2 family), soon SHA-3**

Cryptographic One-Way Hash Functions (2)

- Applications:
 - As a building block in **secure communication protocols** to protect the integrity of messages (typically together with a secret key → MAC)
 - As a building block for **digital signatures**
 - To compute **file integrity checks** to detect data tampering (hash must be stored offline so it cannot easily be tampered with)
 - To store **hashed (and salted) passwords**
- Strengths:
 - **Very fast**, even on small devices feasible
- Limitations:
 - During the recent years, **security vulnerabilities** were found in MD5 and SHA-1 – which is a problem as no widely accepted alternative is available
 - As a result, a **SHA-3 contest** (similar to AES) was started by NIST in 2007, the winning algorithm **Keccak** was announced in October 2012

Message Authentication Codes (MAC) (1)

- Used to protect the **authenticity and integrity** of a message
 - Prerequisite: sender and recipient share a **secret key**

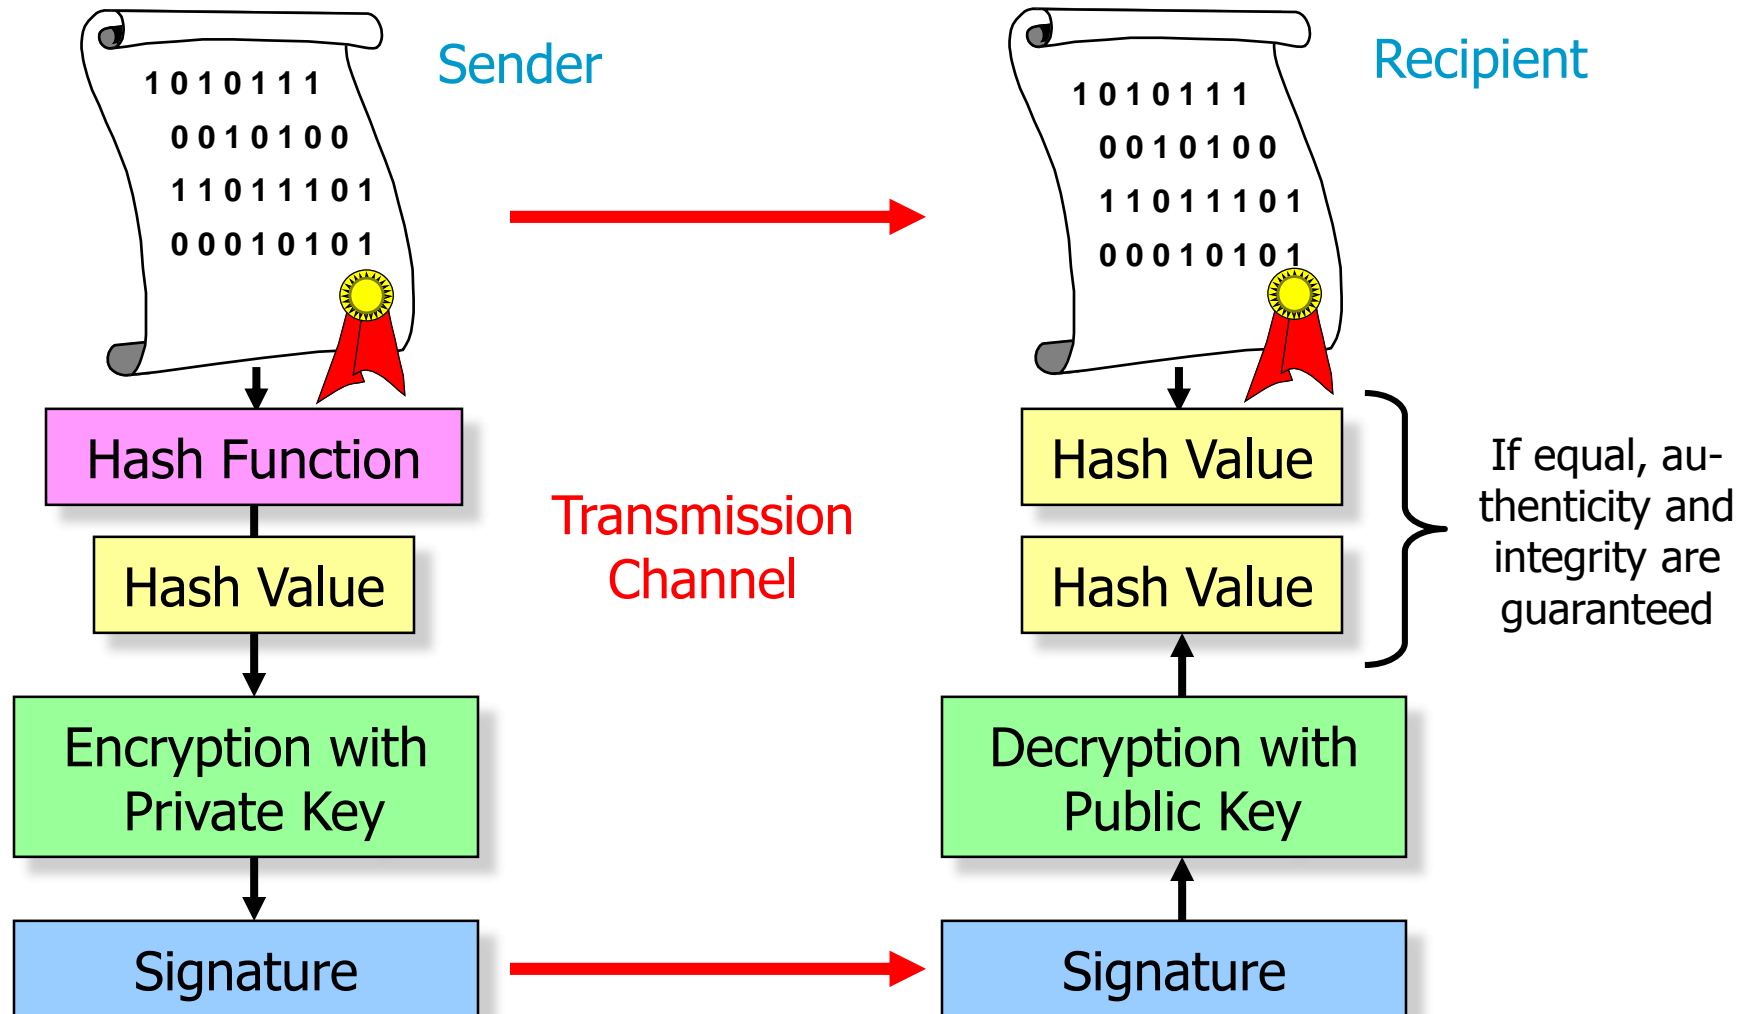


Message Authentication Codes (MAC) (2)

- Standards: **HMAC** (based on keyed hash function), **CBC-MAC** protocol (based on secret key cipher)
- Applications:
 - As a building block in **secure communication protocols** to protect the authenticity and integrity of messages
 - To compute file integrity checks to detect data tampering (key must be stored offline)
- Strengths:
 - **Very fast**, even on small devices feasible
- Limitations:
 - **Often based on MD5 and SHA-1** and therefore relying on partly broken algorithms

Digital Signatures (1)

- Used to guarantee the **authenticity and integrity** of a message
 - Prerequisite: sender has a key pair and recipient knows the public key



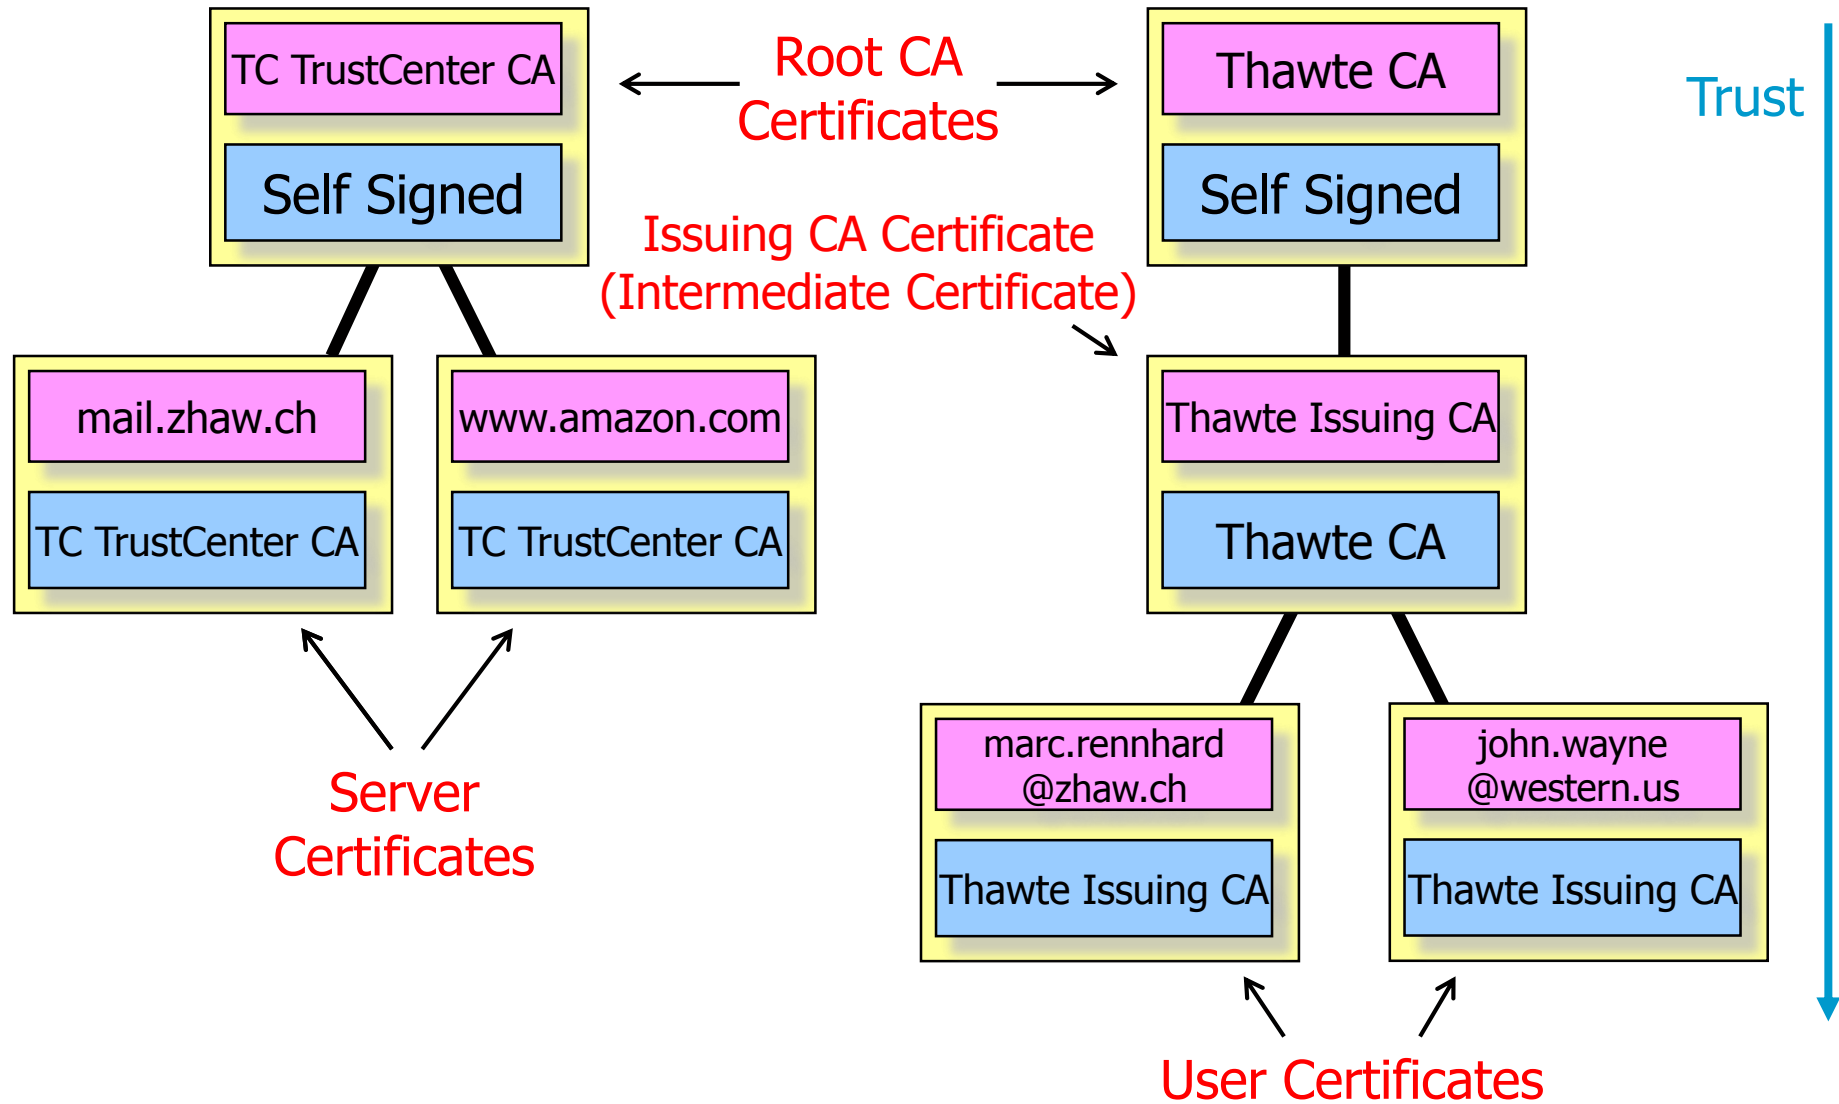
Digital Signatures (2)

- Standards: **RSA** (dominating), DSA, ECDSA
- Applications:
 - As a building block in **secure communication protocols** during initial authentication of the endpoints
 - Digitally signing e-mails to protect the authenticity and integrity
 - Digitally signing of documents or contracts (legally binding)
- Strengths:
 - Enables **strong authentication**
 - Legally binding signatures can significantly increase the efficiency of workflows
- Limitations:
 - Private key must be very well protected as it often has a long lifetime
 - In particular, private key should not be stored in plaintext on user computers
 - Ideally, the private key should be stored on dedicated secure hardware (**smartcards**) → costs, requires software installation
 - Smartcards can usually not be used on mobile devices (e.g. smartphones)

Digital Certificates (1)

- **Digital Certificates** bind a public key to an identity
- Most popular standard: **X.509 certificates** (RFC 2459)
- Certificates are typically issued by **Certification Authorities (CA)**
 - Various official CAs: Entrust, Thawte, VeriSign etc.
 - The CA **checks the identity** of the applicant (e.g. is he/she allowed to get a certificate for domain.com?)
 - A certificate is **digitally signed** by the issuing CA
 - For internal use within e.g. a company, one can also implement an **own internal CA**, which distributes certificates to internal servers and users
- How to **check the validity** of a certificate?
 - Requires that the client/server knows the root CA certificate
 - The public key in this root CA certificate is then used to check the signature on the certificate
 - Root CA certificates are preconfigured in web browsers, e-mail clients etc.

Digital Certificates (2) – Public Key Infrastructure (PKI)



Digital Certificates (3)

- Applications:
 - Basically **whenever public key cryptography is used** and when the public key must be bound to an identity, e.g. e-mail signatures or server authentication in SSL/TLS
- Strengths:
 - **Solves the problem of binding a public key to an identity**, which is a fundamental requirement to enable secure communication
 - Certificate-checking is usually **transparent for the user**, as root CA certificates are preconfigured in certificate-aware software
- Limitations:
 - Correct usage requires **checking for revoked certificates** via OCSP or CRLs, which are sometimes disabled per default
 - The security of many security standards depends on certificates, which makes them **attractive targets**
 - Successful attacks against CAs have happened (compromising CA computers, getting a certificate with a wrong identity...)

User Authentication

Authentication is based on...

- What you **know** (password, PIN, shared secret)
 - Ask for something only the "real" user knows
- What you **have** (Certificate/Private Key, Token, Scratch List, Mobile Phone...)
 - Test for the presence of something only the "real" user has


Username:	<input type="text" value="rennhard"/>
Password:	<input type="text" value="aznHu4Um"/>




01 Z4GH	06 IKDX
02 67TS	07 9PL7
03 UR2A	08 NFLB
04 TIQV	09 K91D
05 3Z5P	10 HA85



- What you **are** (biological pattern, e.g. fingerprint)
 - Non-forgeable biological or behavioural characteristics of the user




Fingerprint



Iris/Retina Scanning



Voice



Face

- **Strong authentication:** combine (at least) two different factors

Passwords (1)

- Very popular, easy to implement and use
 - In particular: no additional client-side soft- or hardware is needed, which is one main reason they are still so widely used
- Basic principle:
 - Every user of the system has a unique user id and a password
 - The server stores all user ids and their passwords in a file or a database
 - To authenticate, the user submits user id and password to the server; the server compares the submitted data to the stored data
 - To increase protection from password theft, passwords should not be stored on the server in plaintext
 - Passwords should be salted and hashed

Username:	<input type="text" value="rennhard"/>
Password:	<input type="password" value="aznHu4Um"/>

Passwords (2)

- Applications:
 - User authentication **at virtually every service**: websites, e-mail servers, local and remote system logins...
- Strengths:
 - **Simple to use** (in the sense that people know how to use them), **cheap** to implement/offer, **portable**
- Limitations:
 - People have problems **picking good passwords** (too short, only characters, based on words of a natural language)
 - Passwords are often **written down, given to others...**
 - Passwords are often **reused** across multiple systems
 - Passwords are susceptible to **phishing attacks**
 - Passwords are (still) sometimes **transmitted across insecure communication links**

One-time Passwords

- Idea:

- Basically similar to passwords, but a one-time password **can only be used once**
- Implementations: scratch lists, electronic tokens (time- or event-based), SMS messages from the service provider

01 Z4GH	06 IKDX
02 67TS	07 9PL7
03 UR2A	08 NFLB
04 TIQV	09 K91D
05 3Z5P	10 HA85



Datum und Uhrzeit letztes
Login: 09.10.2013 20:11
mTAN für neues Login: 147245

- Applications:

- Usually **in combination with a "normal" password** to achieve more secure user authentication (compared to password-only) → **two-factor authentication** (which is regarded as strong authentication)

- Strengths:

- **Simple to use**, "paper-versions" are **relatively cheap, portable**

- Limitations:

- Tokens and SMS messages are relatively **expensive**

User Certificates

- The basic applications / strengths / limitations were already discussed in the section about **digital signatures**
 - Because when using user certificates for user authentication, a digital signature is usually performed
- Just one additional remark with respect to strong authentication
 - If the private key is stored on a smartcard, then this corresponds to (strong) **two-factor authentication** solution
 - Because one must both possess the smartcard (something you **have**)...
 - ...and also know the PIN (something you **know**) that is usually needed to access the smartcard to perform private key operations



Secure Communication Protocols

Secure Sockets Layer / Transport Layer Security (1)


- SSL/TLS provides a **secure communication channel above TCP**
 - Secure means confidentiality, integrity, authenticity
 - There's also a variant that works on top of UDP (Datagram Transport Layer Security)
- The server typically uses **certificates for authentication**
 - Pre-Shared secrets (PSK) cipher suites are also supported
- Client authentication can also use certificates, but client-authentication is **often not done as part of SSL/TLS**
- Several versions:
 - **SSL 2.0**, Netscape, 1995: contains security flaws, don't use it
 - **SSL 3.0**, Netscape, 1996: fixes some security issues of 2.0
 - **TLS 1.0**, IETF, 1999, RFC 2246: similar to SSL 3.0, new MAC computation
 - **TLS 1.1**, IETF, 2006, RFC 4346: similar to TLS 1.0, some security fixes
 - **TLS 1.2**, IETF, 2008, RFC 5246, similar to TLS 1.1, some security fixes

Secure Sockets Layer / Transport Layer Security (2)

- Applications:
 - To secure many TCP-based application protocols: [https](#), [pop3s](#), [imaps](#), [smtps](#)...
 - As a basis for [OpenVPN](#) to build user-space virtual private networks
- Strengths:
 - [Well-established](#) and analyzed
 - Runs in the [user space](#) and can therefore be easily integrated in application software (web, e-mail,...)
- Limitations:
 - Some [attacks](#) have been published in the past years
 - Renegotiation attack, BEAST, CRIME, BREACH, Lucky 13 (all are difficult to carry out in practice)
 - TLS 1.1/1.2 fix some of them, but servers mostly support only TLS 1.0
 - All modern browser now support TLS 1.2, but sometimes disabled per default

IPsec (1)

- IPsec provides a **secure communication channel above IP**
 - Secure means confidentiality, integrity, authenticity
- Several **authentication methods** are supported
 - In practice, certificates and PSK are dominant
- Can be used end-to-end (**transport mode**) or for VPNs (**tunnel mode**), tunnel mode is clearly dominant



IPsec (2)

- Applications:
 - To provide **VPNs** between remote networks
 - To provide **secure access to networks for mobile users**
- Strengths:
 - **Well-established** and analyzed
 - Supported by a **wide spectrum of systems**, including mobile platforms (smartphones)
 - Protection includes transport layer header and – in tunnel mode – internal IP addresses
- Limitations:
 - **Over-engineered** protocol, much more complex than it would have to be
 - Runs in the **kernel space** and therefore require support by the OS
 - **OpenVPN** and other, proprietary VPN solutions are slowly replacing IPsec

Authorization

Access Control

1. Identification


I am: "Username"

Username: rennhard

Password: *****

2. Authentication

Prove that I am "Username"



3. Authorization

What the user is allowed to do

4. Accounting

What the user has done



Identification +
Authentication +
Authorization +
Accounting =
Access Control

Authorization (1)

- Authorization is primarily used by operating systems and applications to determine what a particular user is allowed to do
- There exist three models for authorization:
 - Discretionary Access Control (DAC)
 - Mandatory Access Control (MAC)
 - Role-Based Access Control (RBAC)

Authorization (2)

- Discretionary Access Control:
 - Access to an object is controlled by the **owner** of an object
 - The owner of the resource (e.g. a file) controls **which subjects** (user, group...) can have access to it **and to what degree**
 - Dominating in all major operating systems, usually implemented with **Access Control Lists** (ACLs)
- Mandatory Access Control:
 - Access is controlled (mandated) by the system, a **system-wide policy** determines who is allowed to do what on the system
 - The policy is configured by a (security) policy administrator
 - In today's operating systems, MAC is sometimes used **in addition to DAC** to guarantee some access restrictions in any case
 - When an object is accessed, first enforce MAC, then apply DAC
 - Linux: SELinux, AppArmor
 - Since Windows Vista: Windows Mandatory Integrity Control

Authorization (3)

- Limitations of DAC and MAC
 - Both are very technical and focused to protect access to low-level objects and as a result, they are mainly used in operating systems
 - For applications, this is less well suited, as access rights in applications are usually concerned with user groups that are allowed to do transactions
 - User groups and transactions in an e-shop:
 - Anonymous users are allowed to browse goods
 - Registered customers are allowed to buy goods
 - Marketing personnel may define discounts
 - For such scenarios, Role-Based Access Control is usually better suited
- Role-Based Access Control:
 - The rights of a user depends on the roles that are assigned to him
 - Typical approach works as follows:
 - First define suited roles that contain the corresponding rights (transactions)
 - Roles in an e-shop: customer, sales, marketing, administrator...
 - Assign the roles to users, which grants the user the rights of his roles
 - During runtime, enforce that a user has only the rights according to his roles

Firewalls

Packet Filtering Firewalls (1)

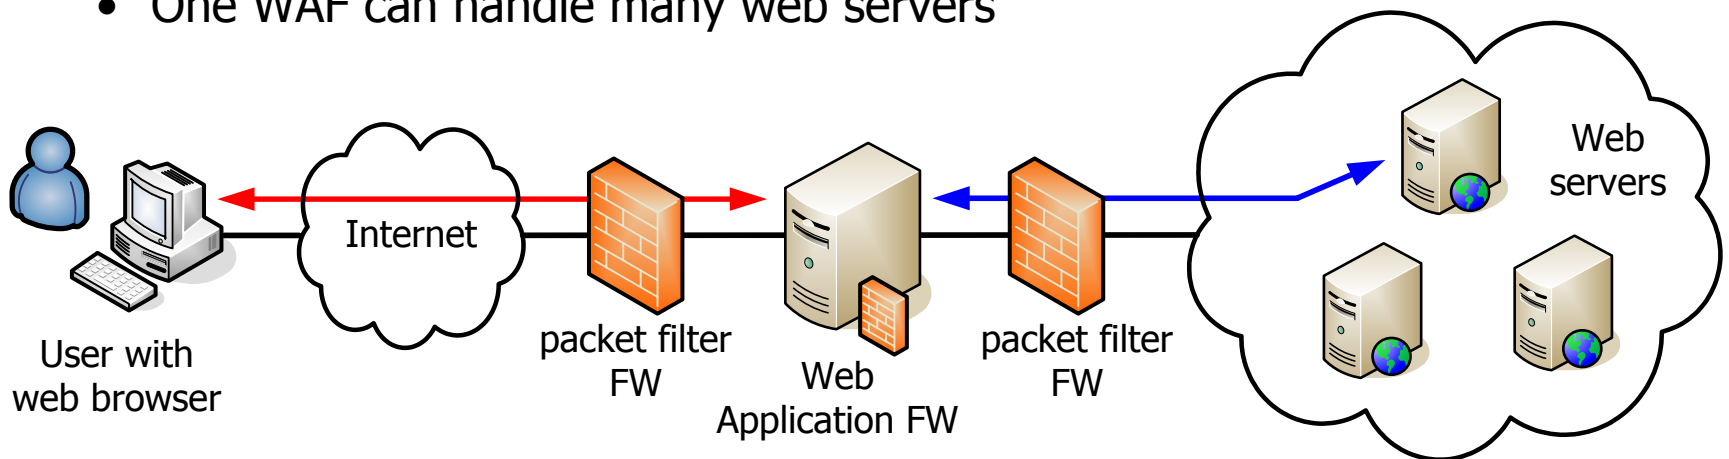
- A **packet filtering firewall** sits between two or more networks to control the packet flow between them
 - The only possibility for packets to travel from one network to another is through the firewall (which may be replicated, though)
- They usually inspect the data in the **network and transport protocol** headers
- Depending on configured **rules**, the traffic is forwarded or blocked
- A typical **rule** could be as follows:
 - Allow traffic from any host on the network 160.85.37.0/24 to port 80 on host 160.85.215.20, but block access to other ports

Packet Filtering Firewalls (2)

- Applications:
 - To filter traffic between connected networks
- Strengths:
 - Well-understood, easy to configure, usually relatively static rule set
 - Blocks a lot of unwanted traffic before it enters an environment
 - Provides a centralised point to control the allowed packet flows, which is much simpler than controlling this at every individual host
- Limitations:
 - Can only control which systems are allowed to communicate, but not the content of the communication

Web Application Firewalls (1)

- Web application firewalls (WAF) **filter traffic at the application layer** and are optimized to protect web applications
- Most frequently, a WAF is operated as a **reverse proxy**
 - Additional packet filtering firewalls guarantee that web servers cannot be reached directly
 - The web browser connect to the WAF, which **relays the traffic** to the desired web server
 - Depending on the configured rules, the WAF decides whether to forward traffic or not
 - One WAF can handle many web servers




Web Application Firewalls (2) – Functionality

- **Content filtering and adaptation**

- Content filtering: e.g. **prevent access** if a request contains possible attack data (Javascript code, SQL statements...)
- Content adaptation: e.g. prevent **information leakage** by replacing replies that contain detailed error messages with a generic error message

- **SSL/TLS termination**: the protected SSL/TLS channel is terminated at the WAF

- This is a prerequisite to allow access to non-encrypted data
- This enables **server certificate management** at a centralised place (and not on every individual web server)



Web Application Firewalls (3) – Functionality

- URL encryption

- The WAF encrypts the paths of all URLs in the HTML page are encrypted with a session-specific key before sending it to the browser
 - So `http://www.mybank.com/actions/paybill` may get to...
 - `http://www.mybank.com/i8RbsO25nwo7+yZqoUjdEW4ljIRQmdo9Rp2Kit`
 - When receiving an encrypted URL from the browser, the WAF decrypts it and forwards it in plain text to the web server
- This means an attacker cannot “produce” valid URLs beyond the ones he receives in the web page – because he does not know the key
 - This protects from forceful browsing and CSRF

- Form protection

- The WAF remembers values of hidden fields sent to the user and compares them with the received values when the form is submitted later
- This prevents attacks where an attacker manipulates hidden fields

- Centralized authentication instead by each web application individually (allows e.g. Single Sign On)

Web Application Firewalls (4)

- Applications:
 - To **further increase protection** for critical web applications (defense in depth, e.g. with e-banking)
 - To provide **protection for insecure web application** that you cannot (or don't want to) fix
 - To provide a **baseline security** for several web applications
 - To provide **external patching** (or just-in-time-patching) to overcome the time between security flaw discovery and application patching
- Strengths:
 - **Powerful security device** if configured correctly
- Limitations:
 - Must be **configured and fine-tuned** to work well in a specific environment (to get low rates of both false positives and negatives)
 - Rule set must be **regularly updated** to detect new attacks

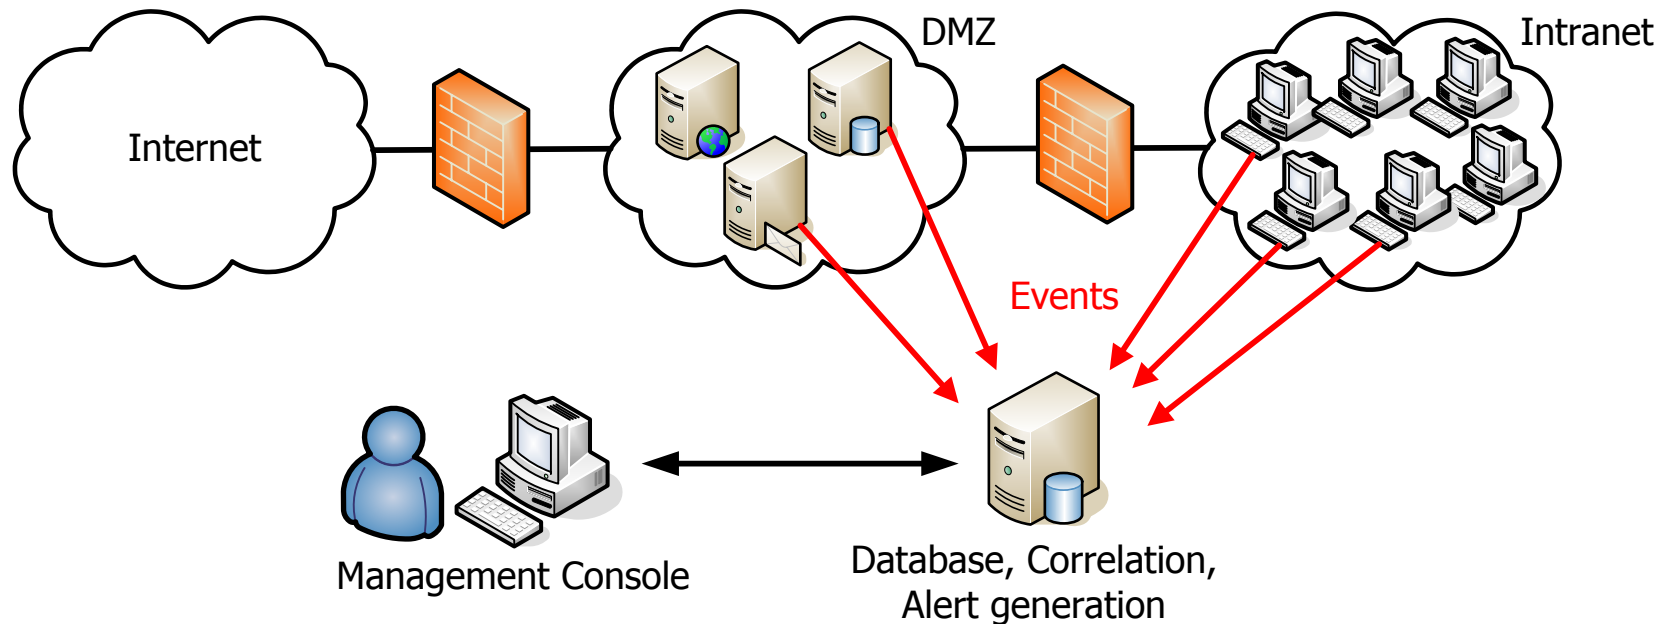
Intrusion Detection / Prevention Systems

Intrusion Detection Systems (1)

- Intrusion Detection Systems (IDS) **monitor network and/or system activity** to detect attacks, attack attempts and general abuse
- Types of IDS:
 - **Host-based IDS** (HIDS) monitoring individual hosts
 - **Network IDS** (NIDS) monitoring entire (or parts of) networks
 - Often used in combination: **Hybrid IDS**
- In general, IDS consist of **multiple components**:
 - Host-based and network sensors, centralised database, correlation engine, management console
 - In the case of a small IDS (e.g. to protect just a single host), these components can all be located on the same system

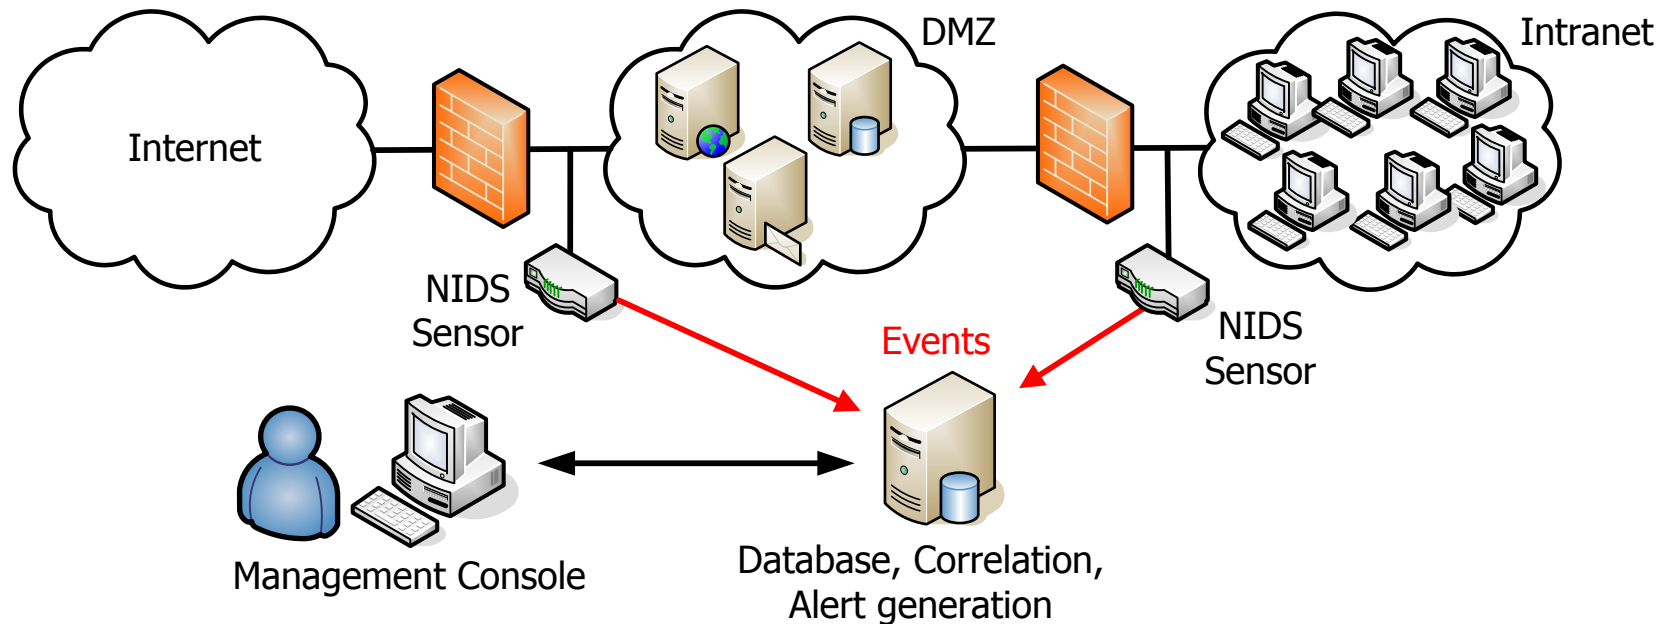
Intrusion Detection Systems (2): Host-based IDS

- Monitor individual **hosts**
- Sensors (software component) are directly **installed on the hosts** that are monitored
- Analyses system calls, application logs, file modifications; but also network data from/to this host
- Examples: OSSEC, Sentry, Tripwire, Snort (when used on a host)



Intrusion Detection Systems (3): Network IDS

- Monitor **network segments**
- **Sensors are dedicated devices** in the network that see the desired traffic (often attached to the monitor port of a switch)
- Analyses network traffic for malicious activity such as scanning traffic or attack attempts
- Examples: HP TippingPoint, ISS RealSecure, Snort (dedicated sensor)

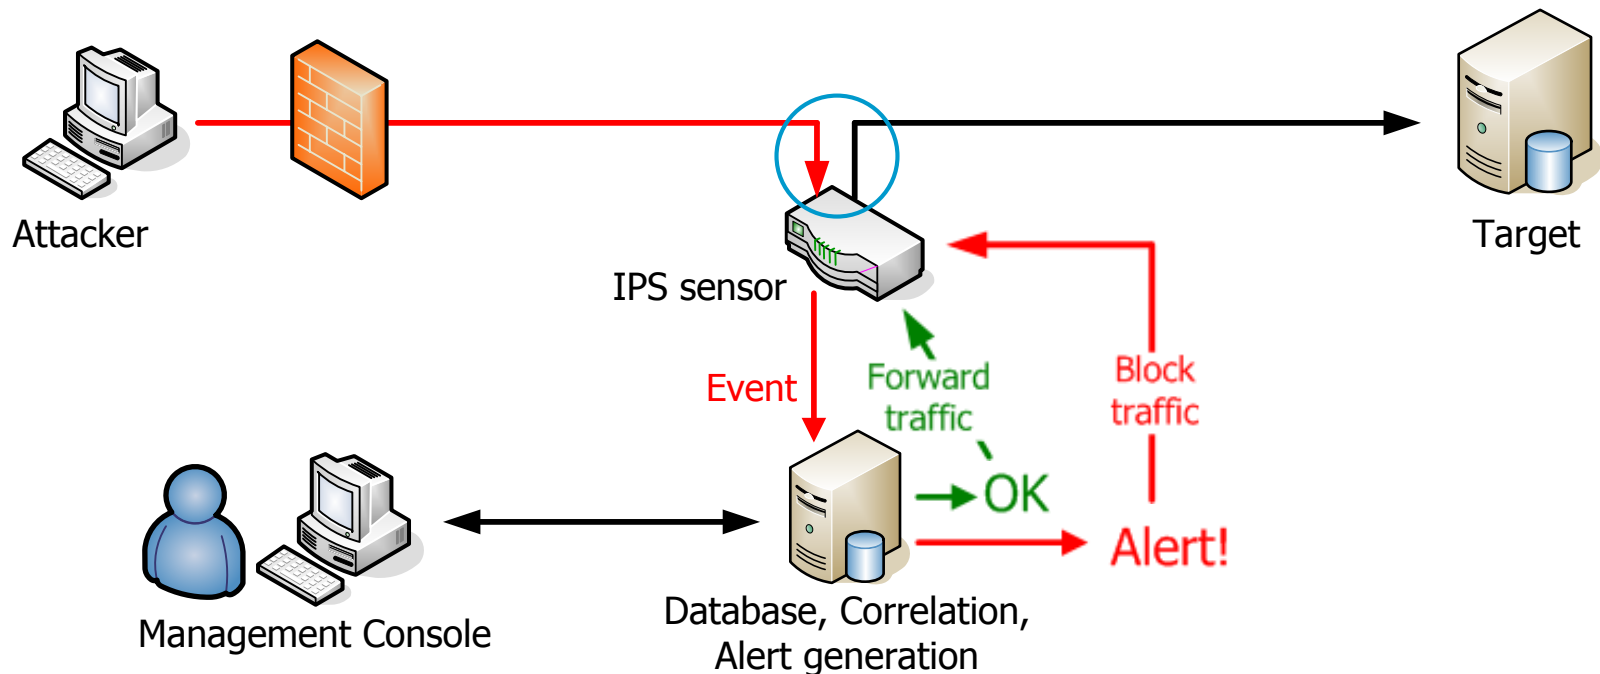


Intrusion Detection Systems (4)

- Applications:
 - To monitor network or system activity to **detect attacks and intrusion attempts**
- Strengths:
 - Provides you with **information** whether and in what way you are under attack (log file analysis can also provide this to a certain degree)
 - Can help **detecting attacks that are very difficult to detect otherwise** (e.g. distributed scans, malware propagation, anomalous system behavior because a system compromise happened...)
- Limitations:
 - Must be **configured and fine-tuned** to work well in a specific environment (low rates of false positives and negatives)
 - Network IDS cannot see through **encrypted traffic** → can be overcome by installing network sensors on hosts
 - An IDS does not stop an attack, it **merely detects and logs** it → human intervention needed (unless the IDS is only used for forensics)

Intrusion Prevention System (IPS) (1)

- Idea: Unlike with IDS, the traffic does not flow past a sensor (passive monitoring), but **through a sensor** (active interception)
 - If a sensor generates an event, the traffic is blocked temporarily
 - If no alert is generated, the traffic is forwarded
 - If the event results in an alert (e.g. by correlating it with other events from other sensors), the **traffic is blocked**
 - Blocked traffic is logged and discarded



Intrusion Prevention Systems (2)

- Advantages compared to IDS:
 - If operated correctly, attacks are not only detected but stopped
 - Because suspicious packets are blocked and only forwarded if no alert is generated
- Additional challenges / limitations compared to IDS:
 - Packets can only be blocked for a very short time to avoid that legitimate communication relationship are significantly delayed
 - This implies that IPS are resource intensive because traffic must be analyzed in near-real time
 - This limits correlation-possibilities with other events from the same or other sensors
 - False positives lock out legitimate user (with IDS, the event/alert is just “registered”)

Summary

- There exist **several security controls** that are helpful when developing secure systems
- This includes:
 - Cryptographic primitives
 - User authentication
 - Secure communication protocols
 - Authorization
 - Firewalls
 - Intrusion detection / prevention systems
 - and more...
- All security controls have their **strengths and limitations**
- When picking a security control, carefully **consider different requirements** (security, usability, costs...) to make good decisions